

%html

<video controls autoplay><source src="https://huggingface.co/ThomasSimonini/ppo-SpaceInvadersNoFrameskip-v4/resolve/main/replay.mp4" type="video/mp4">



!pip install git+https://github.com/DLR-RM/rl-baselines3-zoo



```
Collecting nvidia_cusolver_cu12==11.6.1.9 (from torch<3.0,>=2.3->stable_baselines3<3.0,>=2.5.0->sb3_contrib<3.0,>=2.5.0->rl_zoo3==2.5.0)
  Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia_cusparses-cu12==12.3.1.170 (from torch<3.0,>=2.3->stable_baselines3<3.0,>=2.5.0->sb3_contrib<3.0,>=2.5.0->rl_zoo3==2.5.0)
  Downloading nvidia_cusparses-cu12-12.3.1.170-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in /usr/local/lib/python3.11/dist-packages (from torch<3.0,>=2.3->stable_baselines3<3.0,>=2.5.0->sb3_contrib<3.0,>=2.5.0->rl_zoo3==2.5.0)
Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch<3.0,>=2.3->stable_baselines3<3.0,>=2.5.0->sb3_contrib<3.0,>=2.5.0->rl_zoo3==2.5.0)
Collecting nvidia_nvjitlink-cu12==12.4.127 (from torch<3.0,>=2.3->stable_baselines3<3.0,>=2.5.0->sb3_contrib<3.0,>=2.5.0->rl_zoo3==2.5.0)
  Downloading nvidia_nvjitlink-cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Requirement already satisfied: triton==3.1.0 in /usr/local/lib/python3.11/dist-packages (from torch<3.0,>=2.3->stable_baselines3<3.0,>=2.5.0->sb3_contrib<3.0,>=2.5.0->rl_zoo3==2.5.0)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.11/dist-packages (from torch<3.0,>=2.3->stable_baselines3<3.0,>=2.5.0->sb3_contrib<3.0,>=2.5.0->rl_zoo3==2.5.0)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from sympy==1.13.1->torch<3.0,>=2.3->stable_baselines3<3.0,>=2.5.0->sb3_contrib<3.0,>=2.5.0->rl_zoo3==2.5.0)
Requirement already satisfied: MarkupSafe>=0.9.2 in /usr/local/lib/python3.11/dist-packages (from Mako->alembic>=1.5.0->optuna>=3.0->stable_baselines3<3.0,>=2.5.0->sb3_contrib<3.0,>=2.5.0->rl_zoo3==2.5.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.3.0->stable_baselines3<3.0,>=2.5.0->sb3_contrib<3.0,>=2.5.0->rl_zoo3==2.5.0)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.3.0->stable_baselines3<3.0,>=2.5.0->sb3_contrib<3.0,>=2.5.0->rl_zoo3==2.5.0)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.3.0->stable_baselines3<3.0,>=2.5.0->sb3_contrib<3.0,>=2.5.0->rl_zoo3==2.5.0)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.3.0->stable_baselines3<3.0,>=2.5.0->sb3_contrib<3.0,>=2.5.0->rl_zoo3==2.5.0)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.3.0->stable_baselines3<3.0,>=2.5.0->sb3_contrib<3.0,>=2.5.0->rl_zoo3==2.5.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.3.0->stable_baselines3<3.0,>=2.5.0->sb3_contrib<3.0,>=2.5.0->rl_zoo3==2.5.0)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas>=1.3.0->stable_baselines3<3.0,>=2.5.0->sb3_contrib<3.0,>=2.5.0->rl_zoo3==2.5.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests>=2.28.1->huggingface-hub>=0.16.0->stable_baselines3<3.0,>=2.5.0->sb3_contrib<3.0,>=2.5.0->rl_zoo3==2.5.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests>=2.28.1->huggingface-hub>=0.16.0->stable_baselines3<3.0,>=2.5.0->sb3_contrib<3.0,>=2.5.0->rl_zoo3==2.5.0)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests>=2.28.1->huggingface-hub>=0.16.0->stable_baselines3<3.0,>=2.5.0->sb3_contrib<3.0,>=2.5.0->rl_zoo3==2.5.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests>=2.28.1->huggingface-hub>=0.16.0->stable_baselines3<3.0,>=2.5.0->sb3_contrib<3.0,>=2.5.0->rl_zoo3==2.5.0)
Downloading huggingface_sb3-3.0-py3-none-any.whl (9.7 kB)
Downloading optuna-4.2.1-py3-none-any.whl (383 kB)
383.6/383.6 kB 13.0 MB/s eta 0:00:00
Downloading pytablewriter-1.2.1-py3-none-any.whl (91 kB)
91.1/91.1 kB 7.1 MB/s eta 0:00:00
Downloading sb3_contrib-2.5.0-py3-none-any.whl (92 kB)
92.8/92.8 kB 8.3 MB/s eta 0:00:00
Downloading Shimmy-2.0.0-py3-none-any.whl (30 kB)
Downloading alembic-1.14.1-py3-none-any.whl (233 kB)
233.6/233.6 kB 22.9 MB/s eta 0:00:00
Downloading DataProperty-1.1.0-py3-none-any.whl (27 kB)
Downloading mbstrdecoder-1.1.4-py3-none-any.whl (7.9 kB)
Downloading pathvalidate-3.2.3-py3-none-any.whl (24 kB)
Downloading stable_baselines3-2.5.0-py3-none-any.whl (183 kB)
183.9/183.9 kB 17.4 MB/s eta 0:00:00
Downloading tabledata-1.3.4-py3-none-any.whl (11 kB)
Downloading tcolorpy-0.1.7-py3-none-any.whl (8.1 kB)
Downloading typepy-1.3.4-py3-none-any.whl (31 kB)
Downloading colorlog-6.9.0-py3-none-any.whl (11 kB)
Downloading nvidia_cublas-cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl (363.4 MB)
363.4/363.4 MB 1.4 MB/s eta 0:00:00
Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (13.8 MB)
13.8/13.8 MB 106.2 MB/s eta 0:00:00
Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (24.6 MB)
24.6/24.6 MB 86.7 MB/s eta 0:00:00
Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (883 kB)
883.7/883.7 kB 60.6 MB/s eta 0:00:00
Downloading nvidia_cudnn-cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl (664.8 MB)
664.8/664.8 MB 1.6 MB/s eta 0:00:00
Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl (211.5 MB)
211.5/211.5 MB 5.7 MB/s eta 0:00:00
Downloading nvidia_curand-cu12-10.3.5.147-py3-none-manylinux2014_x86_64.whl (56.3 MB)
56.3/56.3 MB 12.7 MB/s eta 0:00:00
Downloading nvidia_cusolver-cu12-11.6.1.9-py3-none-manylinux2014_x86_64.whl (127.9 MB)
127.9/127.9 MB 8.3 MB/s eta 0:00:00
```

!apt-get install swig cmake ffmpeg



```
Reading package lists... Done
Building dependency tree... Done
```

```

Reading state information... Done
cmake is already the newest version (3.22.1-1ubuntu1.22.04.2).
ffmpeg is already the newest version (7:4.4.2-0ubuntu0.22.04.1).
The following additional packages will be installed:
  swig4.0
Suggested packages:
  swig-doc swig-examples swig4.0-examples swig4.0-doc
The following NEW packages will be installed:
  swig swig4.0
0 upgraded, 2 newly installed, 0 to remove and 20 not upgraded.
Need to get 1,116 kB of archives.
After this operation, 5,542 kB of additional disk space will be used.
Get:1 http://archive.ubuntu.com/ubuntu jammy/universe amd64 swig4.0 amd64 4.0.2-1ubuntu1 [1,110 kB]
Get:2 http://archive.ubuntu.com/ubuntu jammy/universe amd64 swig all 4.0.2-1ubuntu1 [5,632 B]
Fetched 1,116 kB in 1s (769 kB/s)
Selecting previously unselected package swig4.0.
(Reading database ... 124926 files and directories currently installed.)
Preparing to unpack .../swig4.0_4.0.2-1ubuntu1_amd64.deb ...
Unpacking swig4.0 (4.0.2-1ubuntu1) ...
Selecting previously unselected package swig.
Preparing to unpack .../swig_4.0.2-1ubuntu1_all.deb ...
Unpacking swig (4.0.2-1ubuntu1) ...
Setting up swig4.0 (4.0.2-1ubuntu1) ...
Setting up swig (4.0.2-1ubuntu1) ...
Processing triggers for man-db (2.10.2-1) ...

```

To be able to use Atari games in Gymnasium we need to install atari package. And accept-rom-license to download the rom files (games files).

```

!pip install gymnasium[atari]
!pip install gymnasium[accept-rom-license]

```

```

➔ Requirement already satisfied: gymnasium[atari] in /usr/local/lib/python3.11/dist-packages (1.0.0)
Requirement already satisfied: numpy>=1.21.0 in /usr/local/lib/python3.11/dist-packages (from gymnasium[atari]) (1.26.4)
Requirement already satisfied: cloudpickle>=1.2.0 in /usr/local/lib/python3.11/dist-packages (from gymnasium[atari]) (3.1.1)
Requirement already satisfied: typing-extensions>=4.3.0 in /usr/local/lib/python3.11/dist-packages (from gymnasium[atari]) (4.12.2)
Requirement already satisfied: farama-notifications>=0.0.1 in /usr/local/lib/python3.11/dist-packages (from gymnasium[atari]) (0.0.4)
Requirement already satisfied: ale-py>=0.9 in /usr/local/lib/python3.11/dist-packages (from gymnasium[atari]) (0.10.1)
Requirement already satisfied: gymnasium[accept-rom-license] in /usr/local/lib/python3.11/dist-packages (1.0.0)
WARNING: gymnasium 1.0.0 does not provide the extra 'accept-rom-license'
Requirement already satisfied: numpy>=1.21.0 in /usr/local/lib/python3.11/dist-packages (from gymnasium[accept-rom-license]) (1.26.4)
Requirement already satisfied: cloudpickle>=1.2.0 in /usr/local/lib/python3.11/dist-packages (from gymnasium[accept-rom-license]) (3.1.1)
Requirement already satisfied: typing-extensions>=4.3.0 in /usr/local/lib/python3.11/dist-packages (from gymnasium[accept-rom-license])
Requirement already satisfied: farama-notifications>=0.0.1 in /usr/local/lib/python3.11/dist-packages (from gymnasium[accept-rom-license])

```

```

%capture
!apt install python-opengl
!apt install xvfb
!pip3 install pyvirtualdisplay

```

```

# Virtual display
from pyvirtualdisplay import Display

```

```

virtual_display = Display(visible=0, size=(1400, 900))
virtual_display.start()

```

```

➔ <pyvirtualdisplay.display.Display at 0x7a06ba974a50>

```

```

!python -m rl_zoo3.train --algo _____ --env SpaceInvadersNoFrameskip-v4 -f _____ -c _____

```

```

➔ 2025-02-16 20:59:08.734055: E external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:477] Unable to register cuFFT factory: Attempting
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
E0000 00:00:1739739548.996730 5290 cuda_dnn.cc:8310] Unable to register cuDNN factory: Attempting to register factory for plugin cuDN
E0000 00:00:1739739549.061884 5290 cuda_blas.cc:1418] Unable to register cuBLAS factory: Attempting to register factory for plugin cu
2025-02-16 20:59:09.595380: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CP
To enable the following instructions: AVX2 AVX512F FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
usage: train.py [-h] [--algo {a2c,ddpg,dqn,ppo,sac,td3,ars,crossq,qrdqn,tqc,trpo,ppo_lstm}]
                [--env ENV] [--tb TENSORBOARD_LOG] [--i TRAINED_AGENT]
                [--truncate-last-trajectory TRUNCATE_LAST_TRAJECTORY] [--n N_TIMESTEPS]
                [--num-threads NUM_THREADS] [--log-interval LOG_INTERVAL] [--eval-freq EVAL_FREQ]
                [--optimization-log-path OPTIMIZATION_LOG_PATH] [--eval-episodes EVAL_EPISODES]
                [--n-eval-envs N_EVAL_ENVS] [--save-freq SAVE_FREQ] [--save-replay-buffer]
                [-f LOG_FOLDER] [--seed SEED] [--vec-env {dummy,subproc}] [--device DEVICE]
                [--n-trials N_TRIALS] [--max-total-trials MAX_TOTAL_TRIALS] [--optimize]
                [--no-optim-plots] [--n-jobs N_JOBS] [--sampler {random,tpe,skopt}]
                [--pruner {halving,median,none}] [--n-startup-trials N_STARTUP_TRIALS]
                [--n-evaluations N_EVALUATIONS] [--storage STORAGE] [--study-name STUDY_NAME]

```

```

[--verbose VERBOSE] [--gym-packages GYM_PACKAGES [GYM_PACKAGES ...]]
[--env-kwarg ENV_KWARG [ENV_KWARGS ...]]
[--eval-env-kwarg EVAL_ENV_KWARG [EVAL_ENV_KWARGS ...]]
[-params HYPERPARAMS [HYPERPARAMS ...]] [-conf CONF_FILE] [--uuid] [--track]
[--wandb-project-name WANDB_PROJECT_NAME] [--wandb-entity WANDB_ENTITY] [-P]
[-tags WANDB_TAGS [WANDB_TAGS ...]]
train.py: error: argument --algo: invalid choice: '_____' (choose from 'a2c', 'ddpg', 'dqn', 'ppo', 'sac', 'td3', 'ars', 'crossq', 'c

```

▼ Solution

```
!python -m rl_zoo3.train --algo dqn --env SpaceInvadersNoFrameskip-v4 -f logs/ -c atari-dqn.yml
```

[5]	ep_reward_mean	336
	exploration_rate	0.01
	time/	
	episodes	1924
	fps	290
	time_elapsed	1364
	total_timesteps	396283
	train/	
	learning_rate	0.0003
	loss	0.0198
n_updates	74070	

rollout/	
ep_len_mean	3.07e+03
ep_rew_mean	339
exploration_rate	0.01
time/	
episodes	1928
fps	290
time_elapsed	1369
total_timesteps	397470
train/	
learning_rate	0.0003
loss	0.0211
n_updates	74367

rollout/	
ep_len_mean	3.07e+03
ep_rew_mean	343
exploration_rate	0.01
time/	
episodes	1932
fps	290
time_elapsed	1373
total_timesteps	398736
train/	
learning_rate	0.0003
loss	0.0495
n_updates	74683

```
Eval num_timesteps=400000, episode_reward=348.00 +/- 59.38
Episode length: 3118.80 +/- 240.24
```

eval/	
mean_ep_length	3.12e+03
mean_reward	348
rollout/	
exploration_rate	0.01
time/	
total_timesteps	400000
train/	
learning_rate	0.0003
loss	0.0327
n_updates	74999

Saving to logs//dqn/SpaceInvadersNoFrameskip-v4 4

```
!python -m rl_zoo3.enjoy --algo dqn --env SpaceInvadersNoFrameskip-v4 --no-render --n-timesteps 5000 --folder logs/
```

```

2025-02-16 22:01:02.489803: E external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:477] Unable to register cuFFT factory: Attempting
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
E0000 00:00:1739743262.509909 20976 cuda_dnn.cc:8310] Unable to register cuDNN factory: Attempting to register factory for plugin cuDN
E0000 00:00:1739743262.516098 20976 cuda_blas.cc:1418] Unable to register cuBLAS factory: Attempting to register factory for plugin cu
2025-02-16 22:01:02.539170: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CP
To enable the following instructions: AVX2 AVX512F FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
Loading latest experiment, id=4

```

```
Loading logs/dqn/SpaceInvadersNoFrameskip-v4_4/SpaceInvadersNoFrameskip-v4.zip
A.L.E: Arcade Learning Environment (version 0.10.1+unknown)
[Powered by Stella]
Stacking 4 frames
Atari Episode Score: 575.00
Atari Episode Length 4731
Atari Episode Score: 425.00
Atari Episode Length 3315
Atari Episode Score: 410.00
Atari Episode Length 2987
Atari Episode Score: 385.00
Atari Episode Length 2241
Atari Episode Score: 585.00
Atari Episode Length 4625
```

```
from huggingface_hub import notebook_login # To log to our Hugging Face account to be able to upload models to the Hub.
notebook_login()
!git config --global credential.helper store
```

```
!python -m rl_zoo3.push_to_hub --algo dqn --env SpaceInvadersNoFrameskip-v4 --repo-name dqn-SpaceInvadersNoFrameskip-v4 -orga Vanheart -f

2025-02-16 22:02:28.184611: E external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:477] Unable to register cuFFT factory: Attempti
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
E0000 00:00:1739743348.204912 21351 cuda_dnn.cc:8310] Unable to register cuDNN factory: Attempting to register factory for plugin c
E0000 00:00:1739743348.211242 21351 cuda_blas.cc:1418] Unable to register cuBLAS factory: Attempting to register factory for plugin
2025-02-16 22:02:28.232306: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available
To enable the following instructions: AVX2 AVX512F FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
Loading latest experiment, id=4
Loading logs/dqn/SpaceInvadersNoFrameskip-v4_4/SpaceInvadersNoFrameskip-v4.zip
A.L.E: Arcade Learning Environment (version 0.10.1+unknown)
[Powered by Stella]
Stacking 4 frames
Wrapping the env in a VecTransposeImage.
Uploading to Vanheart/dqn-SpaceInvadersNoFrameskip-v4, make sure to have the rights
i This function will save, evaluate, generate a video of your agent,
create a model card and push everything to the hub. It might take up to some
minutes if video generation is activated. This is a work in progress: if you
encounter a bug, please open an issue.
Fetching 1 files: 0% 0/1 [00:00<?, ?it/s]
.gitattributes: 100% 1.52k/1.52k [00:00<00:00, 9.74MB/s]
Fetching 1 files: 100% 1/1 [00:00<00:00, 3.21it/s]
Saving model to: hub/dqn-SpaceInvadersNoFrameskip-v4/dqn-SpaceInvadersNoFrameskip-v4
Saving video to /tmp/tmpia9k7nvh/-step-0-to-step-1000.mp4
/usr/local/lib/python3.11/dist-packages/moviepy/config_defaults.py:1: DeprecationWarning: invalid escape sequence '\P'
"""
Moviepy - Building video /tmp/tmpia9k7nvh/-step-0-to-step-1000.mp4.
Moviepy - Writing video /tmp/tmpia9k7nvh/-step-0-to-step-1000.mp4

Moviepy - Done !
Moviepy - video ready /tmp/tmpia9k7nvh/-step-0-to-step-1000.mp4
X 'DummyVecEnv' object has no attribute 'video_recorder'
X We are unable to generate a replay of your agent, the package_to_hub
process continues
X Please open an issue at
https://github.com/huggingface/huggingface\_sb3/issues
i Pushing repo dqn-SpaceInvadersNoFrameskip-v4 to the Hugging Face
Hub
policy.optimizer.pth: 0% 0.00/13.5M [00:00<?, ?B/s]
policy.pth: 0% 0.00/13.5M [00:00<?, ?B/s]

pytorch_variables.pth: 0% 0.00/864 [00:00<?, ?B/s]

dqn-SpaceInvadersNoFrameskip-v4.zip: 0% 0.00/27.2M [00:00<?, ?B/s]

Upload 5 LFS files: 0% 0/5 [00:00<?, ?it/s]

train_eval_metrics.zip: 0% 0.00/17.3k [00:00<?, ?B/s]

dqn-SpaceInvadersNoFrameskip-v4.zip: 11% 2.93M/27.2M [00:00<00:00, 25.5MB/s]
pytorch_variables.pth: 100% 864/864 [00:00<00:00, 5.52kB/s]
train eval metrics.zip: 100% 17.3k/17.3k [00:00<00:00, 102kB/s]
```